

Python for Data Analysis and Scientific Computing

X433.3 (2 semester units in COMPSCI)

Instructor Alexander I. Iliev, Ph.D.

Course Content Outline

- **NumPy 2/3**
 - Array operations
 - Reductions
 - Broadcasting
 - Array: shaping, reshaping, flattening, resizing, dimension changing
 - Data sorting

Midterm / Project proposal due

- **NumPy 3/3**
 - Type casting
 - Masking data
 - Organizing arrays
 - Loading data files
 - Dealing with polynomials
 - Good coding practices

- **Scipy 1/2**
 - What is Scipy?
 - Working with files
 - Algebraic operations
 - The Fast Fourier Transform
 - Signal Processing

HW4

- **Scipy 2/2**
 - Interpolation
 - Statistics
 - Optimization

- **Project**
 - Project Presentation

Final Project

Working with files

- Working with files – **playing sounds:**

PyAudio: for playing and recording sounds with Python

Installation steps for PyAudio on Mac OS X:

Step 1: execute this line in your **terminal**:

```
ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```

Step 2: execute this line in your **terminal**:

```
brew install portaudio
```

Step 3: execute this line in your **Python setup (Pyzo)**:

```
pip install pyaudio
```

Step 4: write your routine to play the audio and save it as a .py module

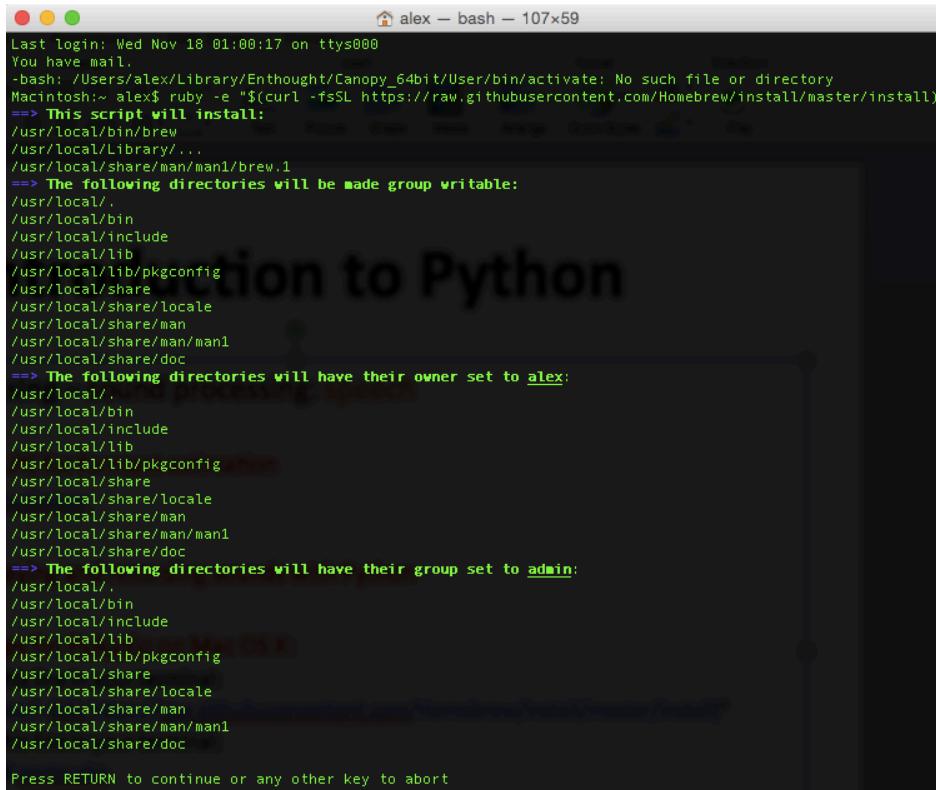
Working with files

- Working with files – playing sounds:

Installation steps for PyAudio on Mac OS X:

Step 1: execute this line in your **terminal**:

```
ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
```



A screenshot of a Mac OS X terminal window titled "alex — bash — 107x59". The window shows the execution of a Homebrew installation script. The output includes messages about file activation, directory permissions, and ownership changes. A large watermark reading "Introduction to Python" is visible across the center of the terminal window.

```
Last login: Wed Nov 18 01:00:17 on ttys000
You have mail.
-bash: /Users/alex/Library/Enthought/Canopy_64bit/User/bin/activate: No such file or directory
Macintosh:~ alex$ ruby -e "$(curl -fsSL https://raw.githubusercontent.com/Homebrew/install/master/install)"
=> This script will install:
/usr/local/bin/brew
/usr/local/Library/...
/usr/local/share/man/man1/brew.1
=> The following directories will be made group writable:
/usr/local/
/usr/local/bin
/usr/local/include
/usr/local/lib
/usr/local/lib/pkgconfig
/usr/local/share
/usr/local/share/locale
/usr/local/share/man
/usr/local/share/man/man1
/usr/local/share/doc
=> The following directories will have their owner set to alex:
/usr/local/
/usr/local/bin
/usr/local/include
/usr/local/lib
/usr/local/lib/pkgconfig
/usr/local/share
/usr/local/share/locale
/usr/local/share/man
/usr/local/share/man/man1
/usr/local/share/doc
=> The following directories will have their group set to admin:
/usr/local/
/usr/local/bin
/usr/local/include
/usr/local/lib
/usr/local/lib/pkgconfig
/usr/local/share
/usr/local/share/locale
/usr/local/share/man
/usr/local/share/man/man1
/usr/local/share/doc
Press RETURN to continue or any other key to abort
```

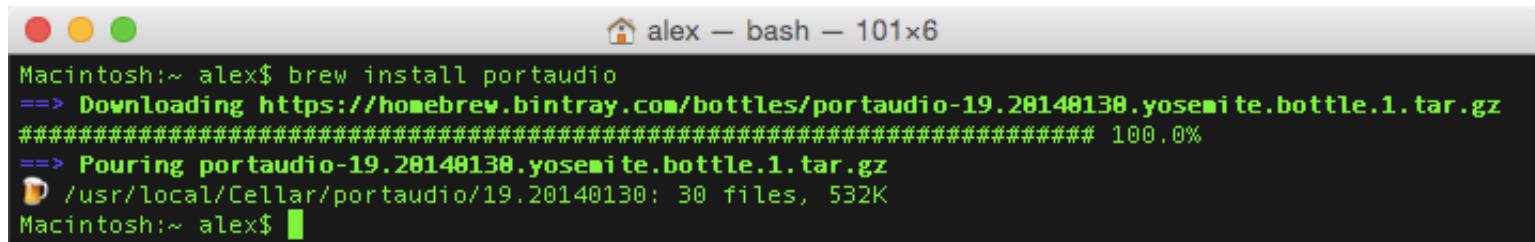
Working with files

- Working with files – playing sounds:

Installation steps for PyAudio on Mac OS X:

Step 2: execute this line in your terminal:

```
brew install portaudio
```



A screenshot of a Mac OS X terminal window titled "alex — bash — 101x6". The window shows the command "brew install portaudio" being run. The output indicates that the file "portaudio-19.20140130.yosemite.bottle.1.tar.gz" is being downloaded from "https://homebrew.bintray.com/bottles/" and then poured into the "/usr/local/Cellar/portaudio/19.20140130" directory. The download progress is shown as a series of hash symbols, reaching 100.0% completion.

```
Macintosh:~ alex$ brew install portaudio
==> Downloading https://homebrew.bintray.com/bottles/portaudio-19.20140130.yosemite.bottle.1.tar.gz
#####
==> Pouring portaudio-19.20140130.yosemite.bottle.1.tar.gz
🍺 /usr/local/Cellar/portaudio/19.20140130: 30 files, 532K
Macintosh:~ alex$
```

Step 3: execute this line in your Python setup (Pyzo):

```
pip install pyaudio
```



A screenshot of the Pyzo Python IDE. The toolbar at the top has a "Python" dropdown. The main area shows the command "In [3]: pip install pyaudio" entered in the terminal. The Pyzo interface includes various toolbars and status indicators.

In [3]: pip install pyaudio

Step 4: write your routine to play the audio and save it as a .py module

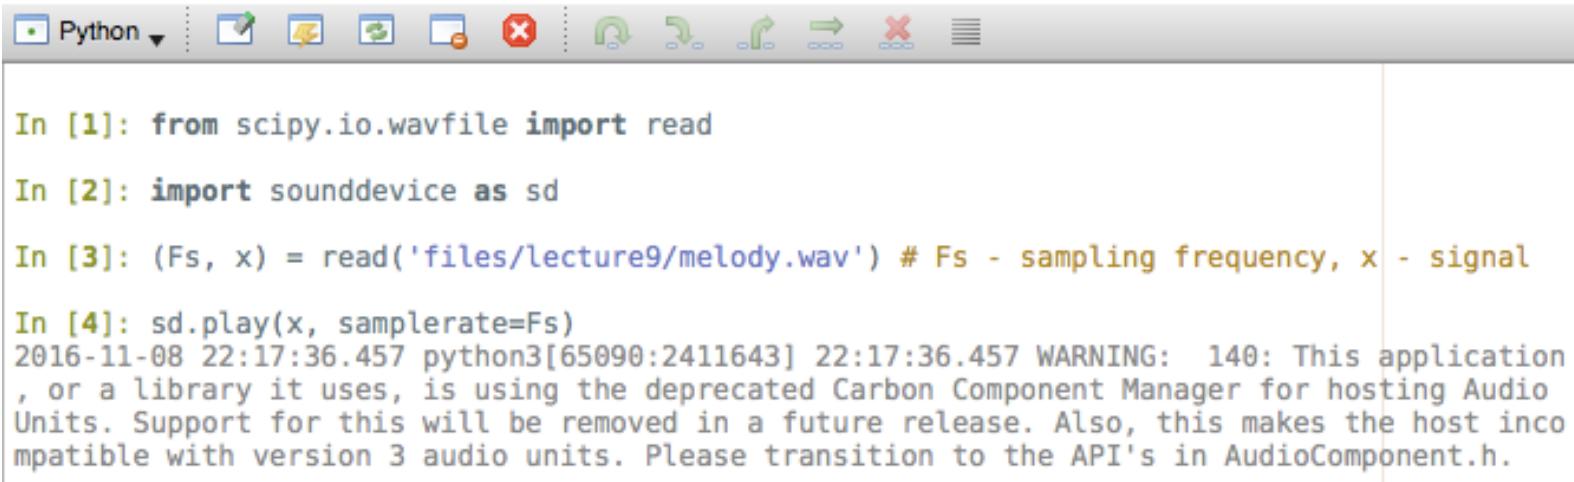
More info for other systems: <http://people.csail.mit.edu/hubert/pyaudio/#binaries>

Working with files

- Working with files – playing sounds with `sounddevice`:

`sounddevice`:

- is **much easier to install and use** as there is no need for you to create your play module
- ... however depending on your setup it might give you a **warning message of deprecated components** so it may not run



The screenshot shows a Jupyter Notebook interface with the following code and output:

```
In [1]: from scipy.io.wavfile import read
In [2]: import sounddevice as sd
In [3]: (Fs, x) = read('files/lecture9/melody.wav') # Fs - sampling frequency, x - signal
In [4]: sd.play(x, samplerate=Fs)
2016-11-08 22:17:36.457 python3[65090:2411643] 22:17:36.457 WARNING: 140: This application
, or a library it uses, is using the deprecated Carbon Component Manager for hosting Audio
Units. Support for this will be removed in a future release. Also, this makes the host inco
mpatible with version 3 audio units. Please transition to the API's in AudioComponent.h.
```

Working with files

- Working with files – recording sounds with sounddevice:

sounddevice:

```
178 ## Working with files – recording sounds with sounddevice 1/2:  
179 from scipy.io.wavfile import read  
180 import sounddevice as sd  
181  
182 # Set the sampling frequency:  
183 Fs = 8000 # kHz  
184  
185 # Duration of our recording in 'sec' will be:  
186 duration = 3  
187  
188 # To record audio data from your sound device into a NumPy array, use:  
189 recorded_audio_1 = sd.rec(int(duration * Fs), samplerate=Fs, channels=2, blocking=True)  
190  
191 # We can now listen to our recording:  
192 sd.play(recorded_audio_1, Fs)
```

Note: make sure you use ‘blocking=True’, to make sure recording has finished

Working with files

- Working with files – recording sounds with sounddevice:

sounddevice:

```
193 ## Working with files – recording sounds with sounddevice 2/2:  
194 from scipy.io.wavfile import read  
195 import sounddevice as sd  
196  
197 # Set the sampling frequency:  
198 Fs = 8000 # kHz  
199  
200 # Duration of our recording in 'sec' will be:  
201 duration = 3  
202  
203 # Let's set our default parameters:  
204 sd.default.samplerate = Fs  
205 sd.default.channels = 2      # for stereo signal  
206  
207 # By default, the recorded array has the data type 'float32', but we can change that:  
208 recorded_audio_2 = sd.rec(duration * Fs, dtype='float64', blocking=True)  
209  
210 # We can now listen to our recording:  
211 sd.play(recorded_audio_2, Fs)
```

Algebraic operations

- Algebraic operations
 - standard **linear algebra operations** in Scipy are provided through the **linalg** module
 - there is a **rich list of prebuilt functions**
 - the input and output of **all linear algebra routines** are **2-dimensional array objects**
 - **scipy.linalg** contains more advanced functions on top of all the functions available in **numpy.linalg**
 - **scipy.linalg** is always **compiled with the ATLAS LAPACK and BLAS libraries**, unlike **numpy.linalg**, hence it has **very fast** linear algebra capabilities
 - **scipy.linalg** is therefore **always better to use** than **numpy.linalg**

Algebraic operations

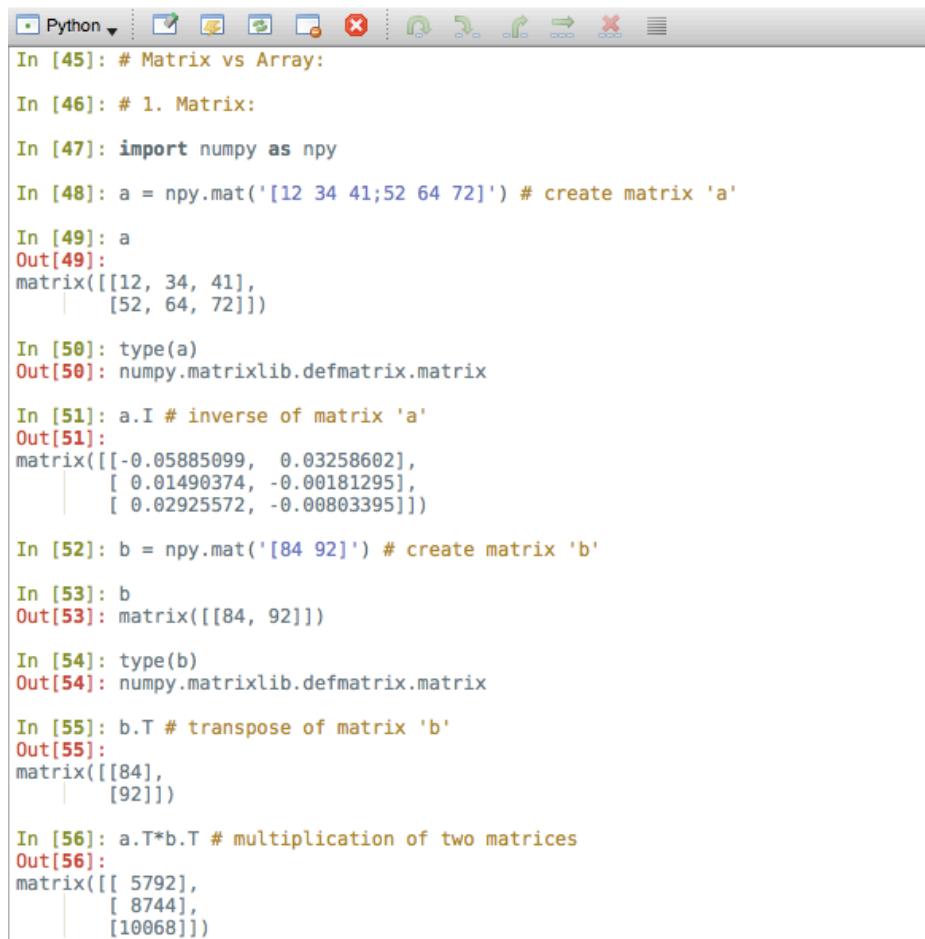
- Algebraic operations
 - Difference between a `numpy.matrix` and 2D `numpy.ndarray`
 - basic operations such as `multiplications` and `transpose` are included in NumPy for both `matrix` and `ndarray` types
 - `numpy.matrix` has a `proper interface` for matrix operations unlike `numpy.ndarray`
 - however, the `numpy.matrix` class `does not add anything` that cannot be achieved by using 2D `numpy.ndarray` objects
 - the implementation of this class `resembles` the one in Matlab
 - `scipy.linalg` operations can be used `just as good` to 2D `numpy.ndarray` objects as well as to `numpy.matrix`

Algebraic operations

- Algebraic operations
 - Difference between a `numpy.matrix` and 2D `numpy.ndarray`

the `I` and `T` class members serve as shortcuts for `inverse` and `transpose` respectively

the `numpy.matrix` class does not add anything that cannot be achieved by using 2D `numpy.ndarray` objects



```
In [45]: # Matrix vs Array:  
In [46]: # 1. Matrix:  
In [47]: import numpy as npy  
In [48]: a = npy.mat('12 34 41;52 64 72') # create matrix 'a'  
In [49]: a  
Out[49]:  
matrix([[12, 34, 41],  
       [52, 64, 72]])  
In [50]: type(a)  
Out[50]: numpy.matrixlib.defmatrix.matrix  
In [51]: a.I # inverse of matrix 'a'  
Out[51]:  
matrix([[-0.05885099,  0.03258602],  
       [ 0.01490374, -0.00181295],  
       [ 0.02925572, -0.000803395]])  
In [52]: b = npy.mat('84 92') # create matrix 'b'  
In [53]: b  
Out[53]: matrix([[84, 92]])  
In [54]: type(b)  
Out[54]: numpy.matrixlib.defmatrix.matrix  
In [55]: b.T # transpose of matrix 'b'  
Out[55]:  
matrix([[84],  
       [92]])  
In [56]: a.T*b.T # multiplication of two matrices  
Out[56]:  
matrix([[ 5792],  
       [ 8744],  
       [10068]])
```

Operations

- Operations
 - Difference between a `numpy.matrix` and 2D `numpy.ndarray`

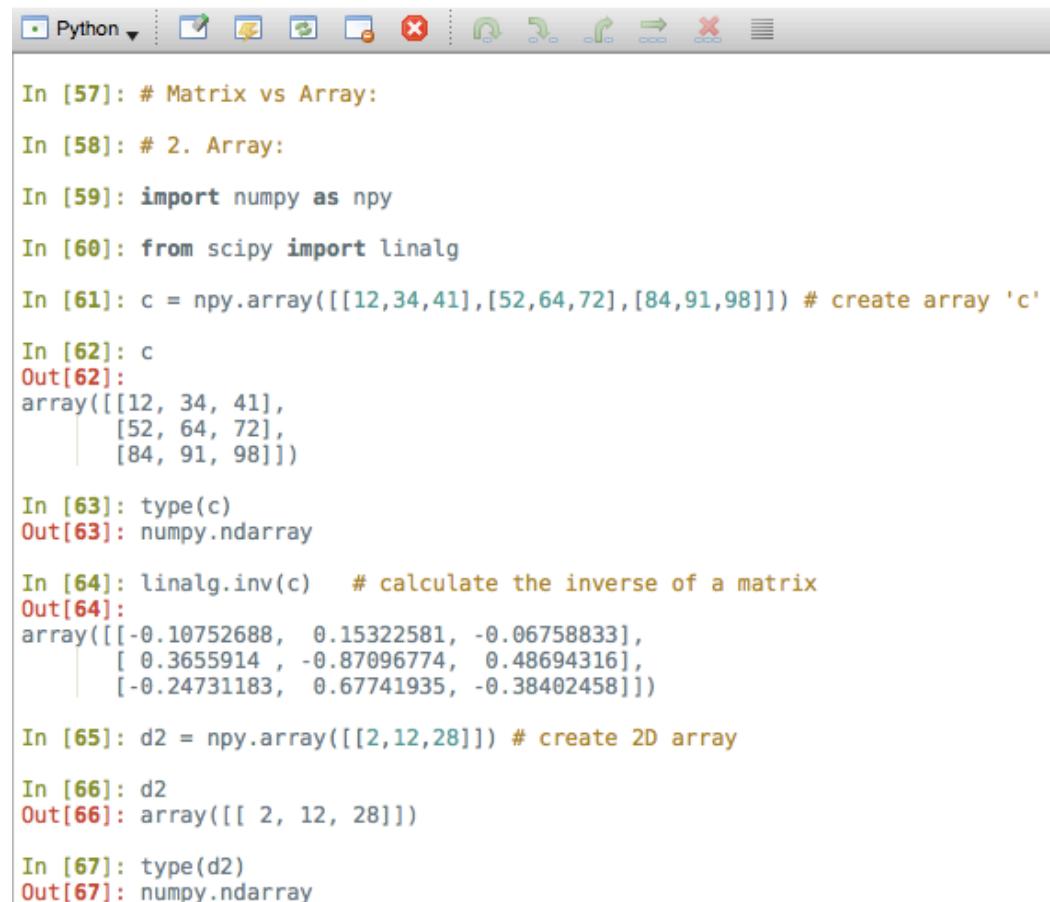
`scipy.linalg` operations can be used just as good to 2D `numpy.ndarray` objects as well as to `numpy.matrix`

Note:

`npy.mat` and `npy.matrix`

Have the same meaning, but are different objects.

Try, using 'id' to see the difference

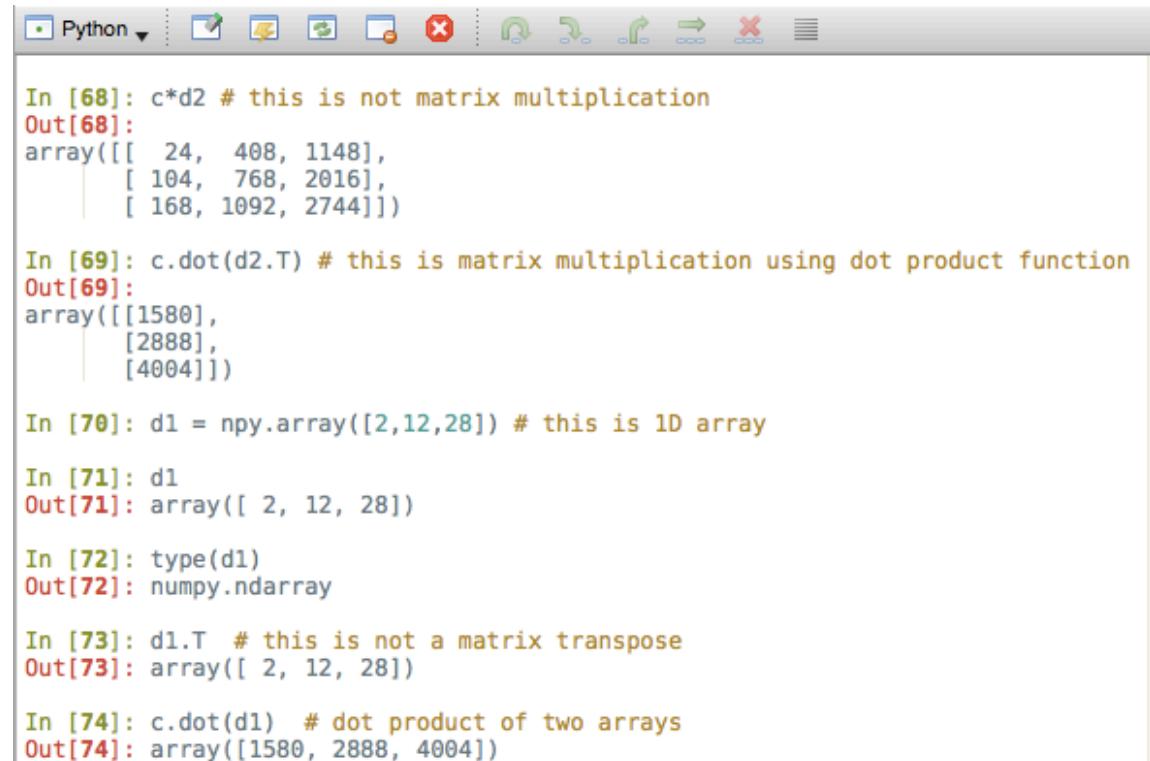


```
In [57]: # Matrix vs Array:  
In [58]: # 2. Array:  
In [59]: import numpy as npy  
In [60]: from scipy import linalg  
In [61]: c = npy.array([[12,34,41],[52,64,72],[84,91,98]]) # create array 'c'  
In [62]: c  
Out[62]:  
array([[12, 34, 41],  
       [52, 64, 72],  
       [84, 91, 98]])  
In [63]: type(c)  
Out[63]: numpy.ndarray  
In [64]: linalg.inv(c) # calculate the inverse of a matrix  
Out[64]:  
array([[-0.10752688,  0.15322581, -0.06758833],  
       [ 0.3655914 , -0.87096774,  0.48694316],  
       [-0.24731183,  0.67741935, -0.38402458]])  
In [65]: d2 = npy.array([[2,12,28]]) # create 2D array  
In [66]: d2  
Out[66]: array([[ 2, 12, 28]])  
In [67]: type(d2)  
Out[67]: numpy.ndarray
```

Operations

- Operations
 - Difference between a `numpy.matrix` and 2D `numpy.ndarray`

`scipy.linalg` operations can be used just as good to 2D `numpy.ndarray` objects as well as to `numpy.matrix`



```
In [68]: c*d2 # this is not matrix multiplication
Out[68]:
array([[ 24,  408, 1148],
       [ 104,  768, 2016],
       [ 168, 1092, 2744]])

In [69]: c.dot(d2.T) # this is matrix multiplication using dot product function
Out[69]:
array([1580,
       2888,
       4004])

In [70]: d1 = np.array([2,12,28]) # this is 1D array

In [71]: d1
Out[71]: array([ 2, 12, 28])

In [72]: type(d1)
Out[72]: numpy.ndarray

In [73]: d1.T # this is not a matrix transpose
Out[73]: array([ 2, 12, 28])

In [74]: c.dot(d1) # dot product of two arrays
Out[74]: array([1580, 2888, 4004])
```

Operations

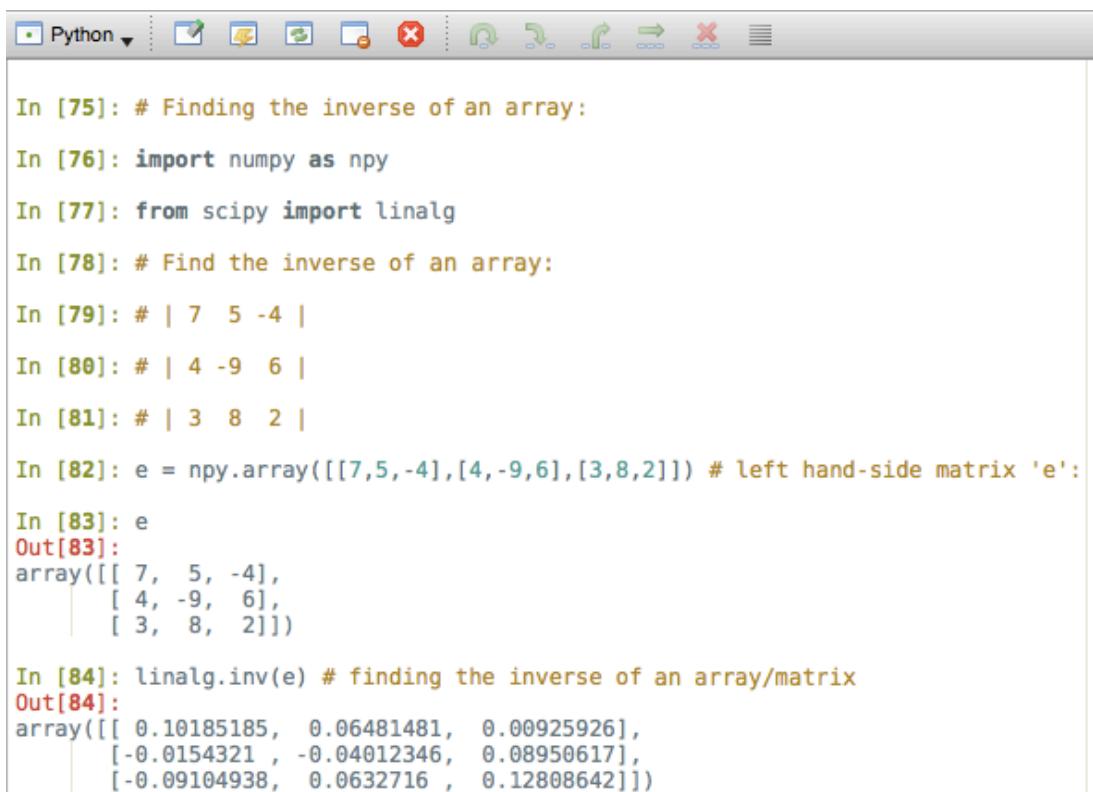
- Operations
 - Finding the inverse of an array (matrix)

matrix 'e' has its **inverse** matrix 'f' such that $e \cdot f = I \rightarrow I$ is the **identity** matrix that has main diagonal with ones

we can then say that: $f = e^{-1}$

the **matrix inverse** of the NumPy array 'e' is obtained in **two ways**:

- using the Scipy `linalg.inv`, or
- using `e.I` when 'e' is a matrix so cast it like this:
`npy.mat(e).I`



```
In [75]: # Finding the inverse of an array:
In [76]: import numpy as npy
In [77]: from scipy import linalg
In [78]: # Find the inverse of an array:
In [79]: # | 7  5 -4 |
In [80]: # | 4 -9  6 |
In [81]: # | 3  8  2 |
In [82]: e = npy.array([[7,5,-4],[4,-9,6],[3,8,2]]) # left hand-side matrix 'e':
In [83]: e
Out[83]:
array([[ 7,  5, -4],
       [ 4, -9,  6],
       [ 3,  8,  2]])

In [84]: linalg.inv(e) # finding the inverse of an array/matrix
Out[84]:
array([[ 0.10185185,  0.06481481,  0.00925926],
       [-0.0154321 , -0.04012346,  0.08950617],
       [-0.09104938,  0.0632716 ,  0.12808642]])
```

Linear Systems

- Linear systems
 - Solving linear systems
 - to compute the **solution vector** of a linear system with **any number of equations** and unknowns, all we have to do is **construct the right hand-side vector** and **left hand-side matrix** with coefficients
 - lets consider the system:
$$\begin{aligned} 7x + 5y - 4z &= 15 \\ 4x - 9y + 6z &= 12 \\ 3x + 8y + 2z &= 10 \end{aligned}$$
 - the **solution vector can be found** by using the **matrix inverse**:
$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} 7 & 5 & -4 \\ 4 & -9 & 6 \\ 3 & 8 & 2 \end{bmatrix}^{-1} \begin{bmatrix} 15 \\ 12 \\ 10 \end{bmatrix} = \begin{bmatrix} 2.39 \\ 0.18 \\ 0.67 \end{bmatrix}$$
 - now lets find the solution using Python
 - we will use **scipy.linalg** as it is faster than **numpy.ndarray**

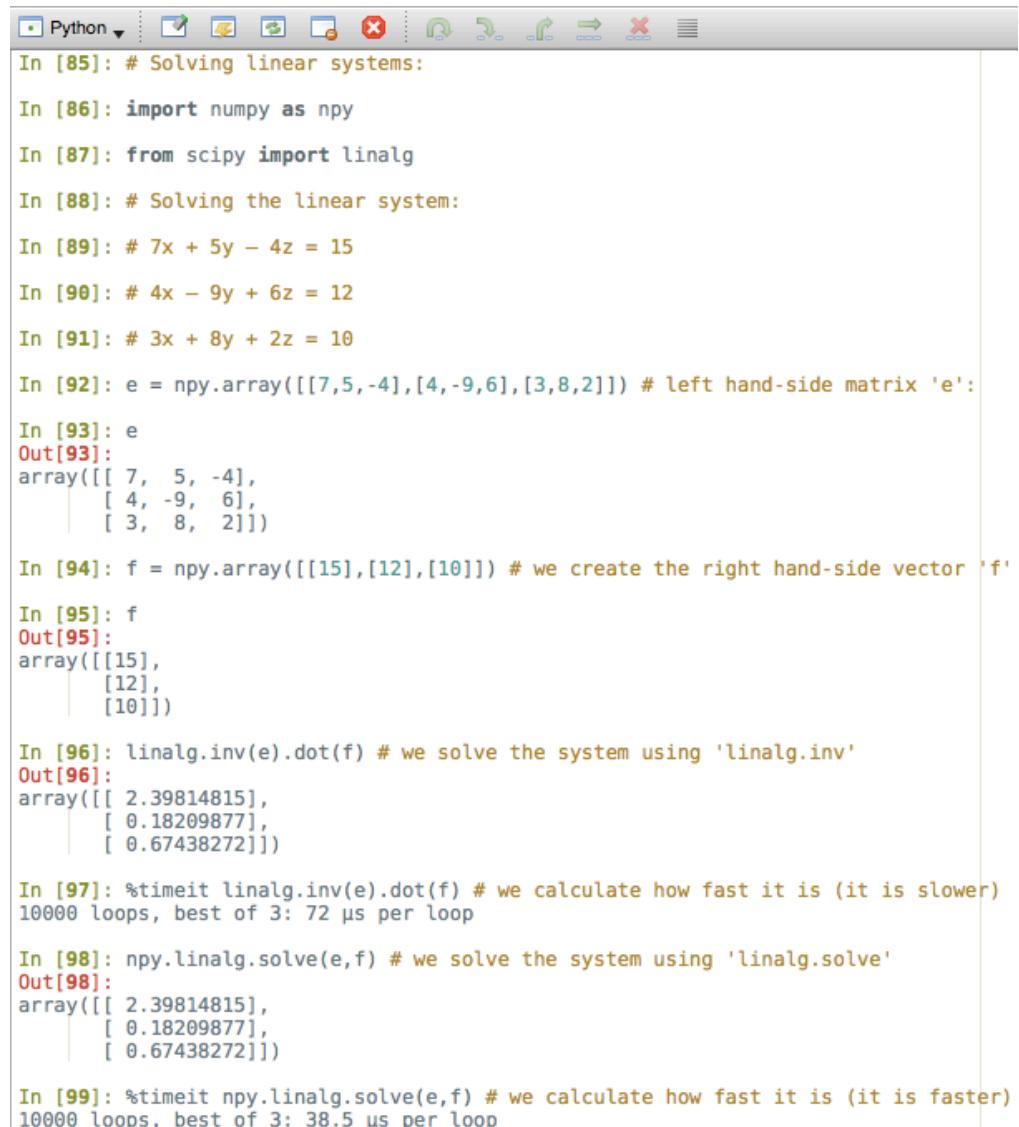
Linear Systems

- Linear systems
 - Solving linear systems

dot product:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^n a_i b_i$$

we will use
`linalg.solve` as it is
faster than `inv.dot`



```
In [85]: # Solving linear systems:
In [86]: import numpy as npy
In [87]: from scipy import linalg
In [88]: # Solving the linear system:
In [89]: # 7x + 5y - 4z = 15
In [90]: # 4x - 9y + 6z = 12
In [91]: # 3x + 8y + 2z = 10
In [92]: e = npy.array([[7,5,-4],[4,-9,6],[3,8,2]]) # left hand-side matrix 'e':
In [93]: e
Out[93]:
array([[ 7,  5, -4],
       [ 4, -9,  6],
       [ 3,  8,  2]])

In [94]: f = npy.array([[15],[12],[10]]) # we create the right hand-side vector 'f'
In [95]: f
Out[95]:
array([[15],
       [12],
       [10]])

In [96]: linalg.inv(e).dot(f) # we solve the system using 'linalg.inv'
Out[96]:
array([[ 2.39814815],
       [ 0.18209877],
       [ 0.67438272]])

In [97]: %timeit linalg.inv(e).dot(f) # we calculate how fast it is (it is slower)
10000 loops, best of 3: 72 µs per loop

In [98]: npy.linalg.solve(e,f) # we solve the system using 'linalg.solve'
Out[98]:
array([[ 2.39814815],
       [ 0.18209877],
       [ 0.67438272]])

In [99]: %timeit npy.linalg.solve(e,f) # we calculate how fast it is (it is faster)
10000 loops, best of 3: 38.5 µs per loop
```

Singular Value Decomposition

- SVD

- Singular Value Decomposition **svd**
- It is a great tool for dimensionality reduction

Can be used in:

- Image compression
- Computer vision (using PCA - SVD)
- Signal processing
- Fitting solutions
- Statistics

```
Python In [100]: # More advanced features:  
In [101]: # Using svd to obtain a Singular Value Decomposition:  
In [102]: e  
Out[102]:  
array([[ 7,  5, -4],  
      [ 4, -9,  6],  
      [ 3,  8,  2]])  
In [103]: a = e + npy.diag([1, 1, 1]) # an array of singular real, non-negative values  
In [104]: a  
Out[104]:  
array([[ 8,  5, -4],  
      [ 4, -8,  6],  
      [ 3,  8,  3]])  
In [105]: u, spectrum, v = linalg.svd(a) # Singular Value Decomposition (a == u*spectrum*v)  
In [106]: spectrum # is the resulting spectrum array  
Out[106]: array([ 13.42424973,   9.1261769 ,   6.28509461])  
In [107]: s = npy.diag(spectrum) # 's' is a matrix with non-zeros in its main diagonal  
In [108]: s  
Out[108]:  
array([[ 13.42424973,   0.          ,   0.          ],  
      [ 0.          ,   9.1261769 ,   0.          ],  
      [ 0.          ,   0.          ,   6.28509461]])  
In [109]: u # ndarray: unitary matrix having left singular vectors as columns  
Out[109]:  
array([[-0.61451975, -0.5313394 , -0.58313285],  
      [ 0.58595351, -0.80234418,  0.1135883 ],  
      [-0.52822719, -0.27188648,  0.80439653]])  
In [110]: v # ndarray: unitary matrix having right singular vectors as rows  
Out[110]:  
array([[-0.30966539, -0.89286512,  0.32695447],  
      [-0.9068147 ,  0.1738915 , -0.38399069],  
      [-0.28599729,  0.41539575,  0.86351139]])  
In [111]: svd_matrix = u.dot(s).dot(v) # to recompose the original matrix  
In [112]: svd_matrix  
Out[112]:  
array([[ 8.,  5., -4.],  
      [ 4., -8.,  6.],  
      [ 3.,  8.,  3.]])
```

Linear algebra operations

- Linear algebra operations
 - there is so much more you can do:
 - determinant of a matrix can be found by using `linalg.det()`
 - computing norms
 - solving linear least-squares problems
 - finding eigenvalues and eigenvectors
 - singular value decomposition SVD (among other decompositions)
 - exponential and logarithm functions
 - trigonometric functions
 - matrix functions and many other special functions

more info: <http://docs.scipy.org/>

The Fast Fourier Transform

- The Fast Fourier Transform – quick intro
 - FFT is the **faster** implementation of the **DFT**
 - FFT is the basis for **frequency analysis** that **converts any signal in time to frequency domain**
 - **fft** is the Fast Fourier Transform (FFT) converts **time**-domain signals **to frequency**-domain
 - **ifft** is the **Inverse** Fast Fourier Transform (IFFT) and converts **frequency to time**-domain
 - the Fourier transform is represented like this:

$$X_k = \sum_{n=0}^{N-1} x_n e^{-i2\pi k \frac{n}{N}}$$

where: $[x_0, \dots, x_{N-1}]$ are complex conjugate numbers and $[k=0, \dots, N-1]$

- the **most commonly** used FFT is the **Cooley–Tukey** algorithm

The Fast Fourier Transform

- The Fast Fourier Transform

Example:

- notice that we only import what we need

- we add three simple tones to create one complex tone

- to go from time domain to frequency domain we use **fft** in two ways

- each frequency bin represents frequency in [Hz] ->

```
243 # FFT example:  
244 from numpy.fft import rfft  
245 from scipy import arange, sin, pi, fft, real ,imag, log10  
246 from scipy.io.wavfile import write  
247 from matplotlib.pyplot import figure, plot, subplot, axis, grid  
248 from pylab import xticks, yticks, xlim, ylim, xlabel, ylabel, title  
249  
250 Fs=4000 # sampling frequency  
251 a = arange(1024) # create a vector holding the number of bins  
252 signal1 = sin(2*pi*a*(650/Fs)) # create a tone with frequency = 650Hz  
253 signal2 = sin(2*pi*a*(1150/Fs)) # create a tone with frequency = 1.15kHz  
254 signal3 = sin(2*pi*a*(1450/Fs)) # create a tone with frequency = 1.425kHz  
255 signal4 = sin(2*pi*a*(1250/Fs)) # create a tone with frequency = 1.25kHz  
256  
257 # Create a complex tone:  
258 signal = signal1 + signal2 + signal3  
259  
260 # Take the FFT of the complex signal:  
261 freq_domain_npy = rfft(signal) # using npy  
262 freq_domain_cpy = fft(signal) # using cpy  
263 bin_val = Fs/len(a) # calculate the value of each frequency bin in [Hz]  
264 bin_val # each bin in [Hz]  
265 t = arange(1,Fs/2+2,bin_val) # create a vector of frequency bins in [Hz]  
266  
267 # Save into a wav file:  
268 write('files/lecture8/fft_file_example.wav',Fs,signal) # save to file
```

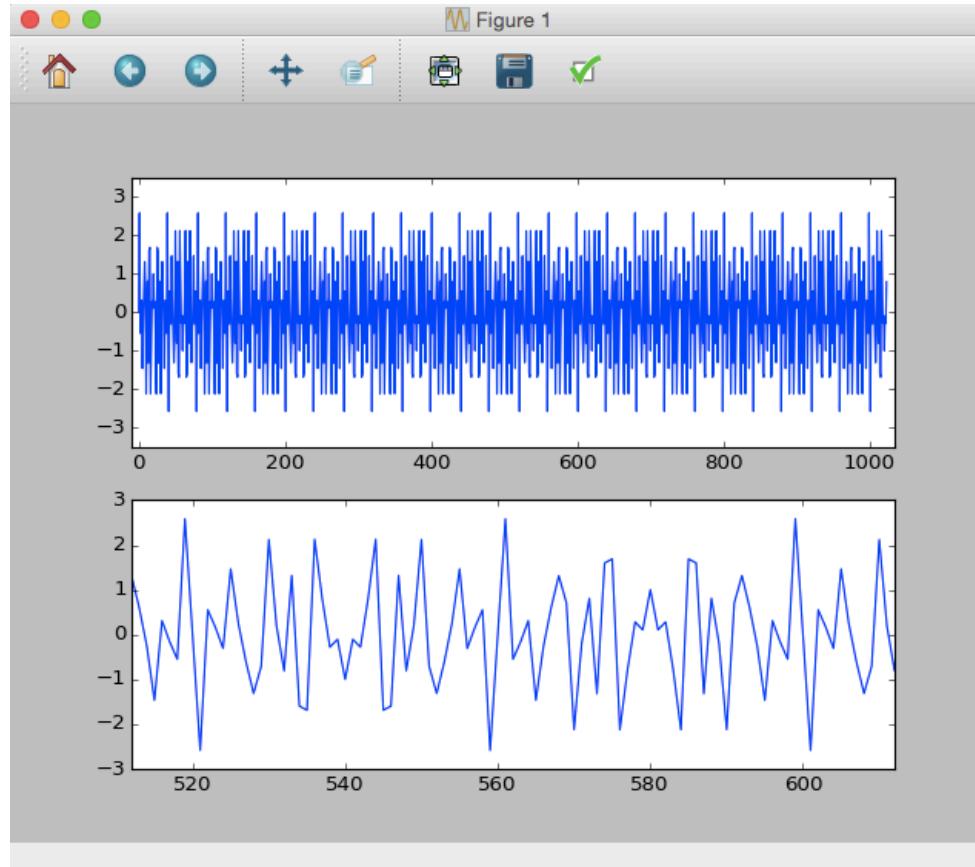


The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
270 # Plot the raw Time domain signal:  
271 figure(1), subplot(2,1,1), plot(signal), xlim(-10, len(a)+10), ylim(-3.5,3.5)  
272 subplot(2,1,2), plot(signal), xlim(len(a)/2,len(a)/2+100) # just a snipped of the signal
```

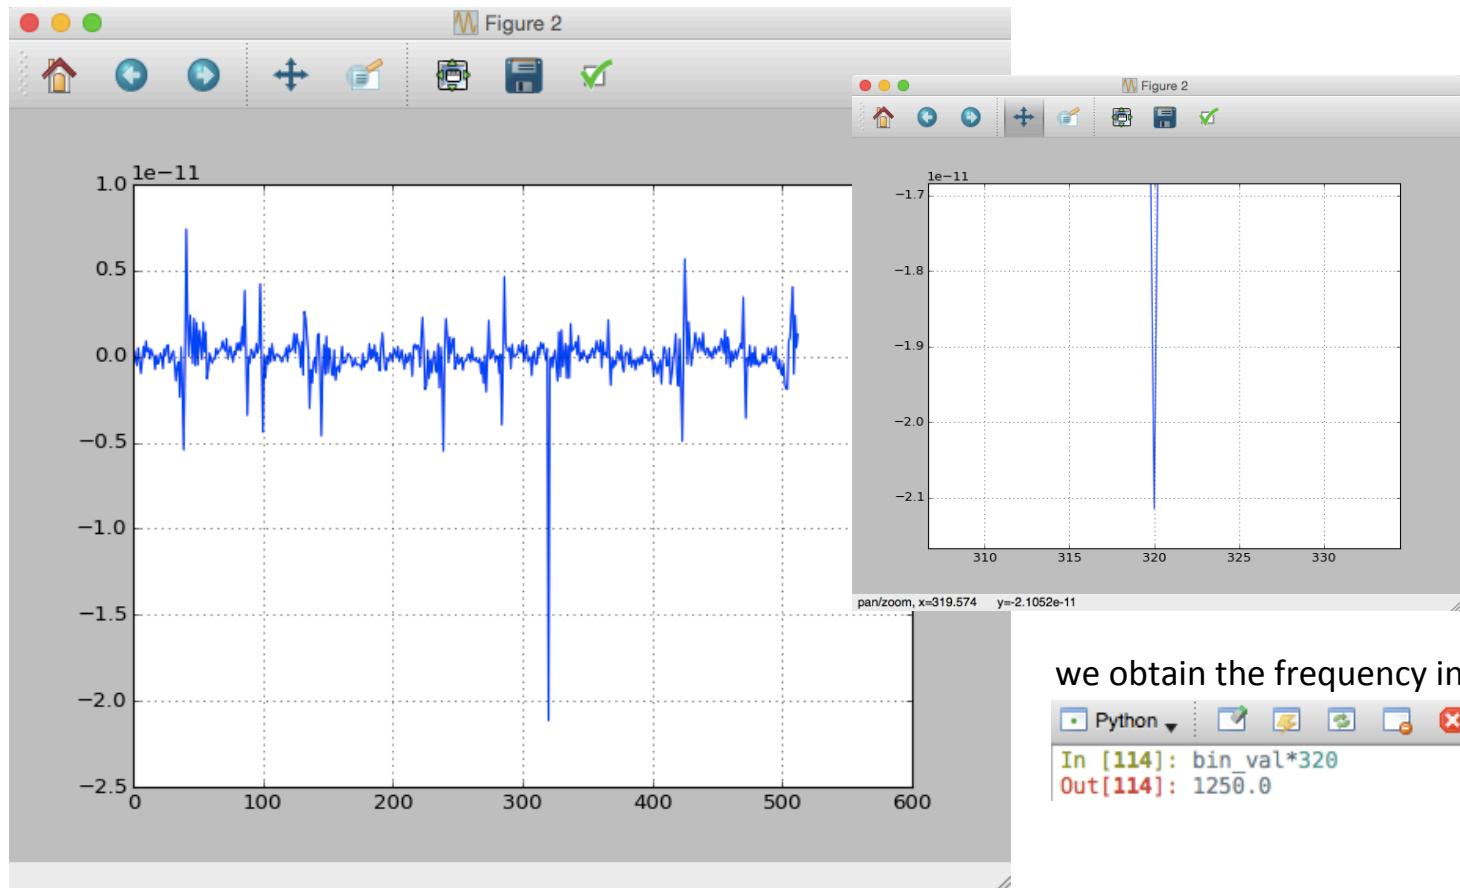


The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
274 # Plot the Frequency domain of 'signal4':  
275 figure(2), plot(rfft(signal4)) # observe the quantization noise due to rounding errors  
276 grid(True)
```



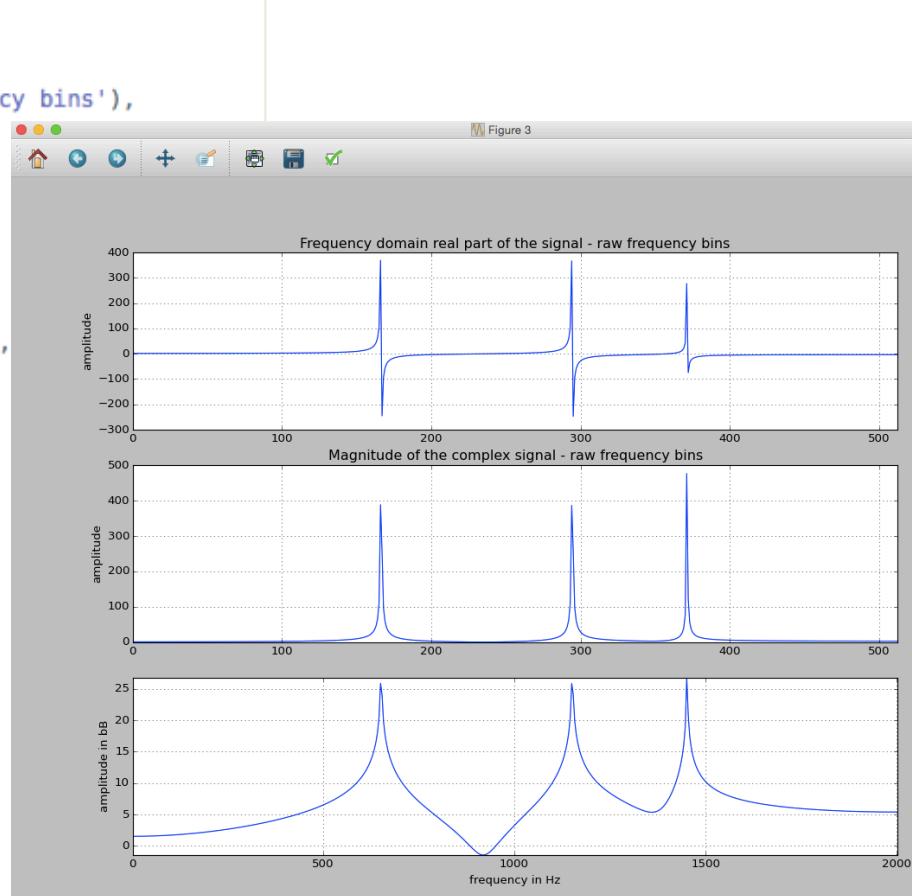
The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
278 # Plot the Frequency domain signal using numpy:  
279 figure(3)  
280 subplot(3,1,1),  
281 title('Frequency domain real part of the signal - raw frequency bins'),  
282 plot(freq_domain_npy),  
283 xlim(0,len(a)/2),  
284 ylabel('amplitude'),  
285 grid(True)  
286  
287 # Plot the magnitude of the complex signal output:  
288 subplot(3,1,2),  
289 plot(abs(freq_domain_cpy)),  
290 title('Magnitude of the complex signal - raw frequency bins'),  
291 xlim(0,len(a)/2),  
292 ylabel('amplitude'),  
293 grid(True)  
294  
295 # Plot in dB scale:  
296 subplot(3,1,3),  
297 plot(t,10*log10(freq_domain_npy)),  
298 xlabel('frequency in Hz'),  
299 ylabel('amplitude in bB'),  
300 axis('tight'),  
grid(True)
```

- notice the difference
in the **x** scales



The Fast Fourier Transform

- The Fast Fourier Transform

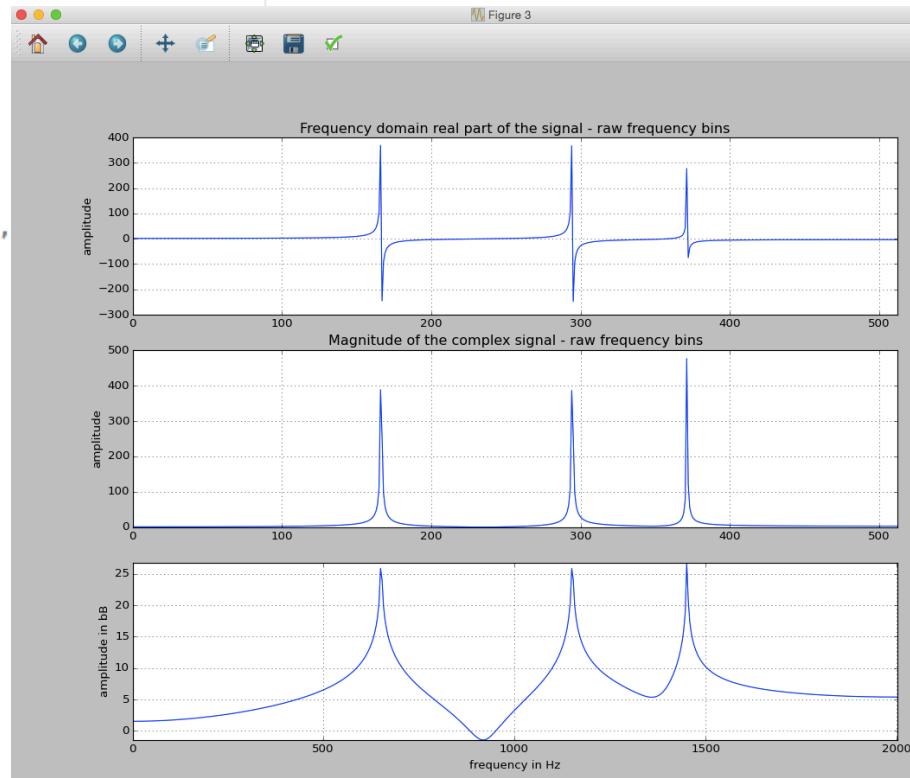
Example:

```
278 # Plot the Frequency domain signal using numpy:  
279 figure(3)  
280 subplot(3,1,1),  
281 title('Frequency domain real part of the signal - raw frequency bins'),  
282 plot(freq_domain_npy),  
283 xlabel('frequency in Hz'),  
284 ylabel('amplitude'),  
285 grid(True)  
286  
287 # Plot the magnitude of the complex signal output:  
288 subplot(3,1,2),  
289 plot(abs(freq_domain_cpy)),  
290 title('Magnitude of the complex signal - raw frequency bins'),  
291 xlabel('frequency in Hz'),  
292 ylabel('amplitude'),  
293 grid(True)  
294  
295 # Plot in dB scale:  
296 subplot(3,1,3),  
297 plot(t,10*log10(freq_domain_npy)),  
298 xlabel('frequency in Hz'),  
299 ylabel('amplitude in bB'),  
300 axis('tight'),  
grid(True)
```

- notice the difference
in the x scales

```
In [11]: bin_val*166  
Out[11]: 648.4375  
In [12]: bin_val*294  
Out[12]: 1148.4375  
In [13]: bin_val*371  
Out[13]: 1449.21875
```

create a vector
 $\begin{aligned} &*(650/F_s) \\ &*(1150/F_s) \\ &*(1450/F_s) \\ &*(1250/F_s) \end{aligned}$



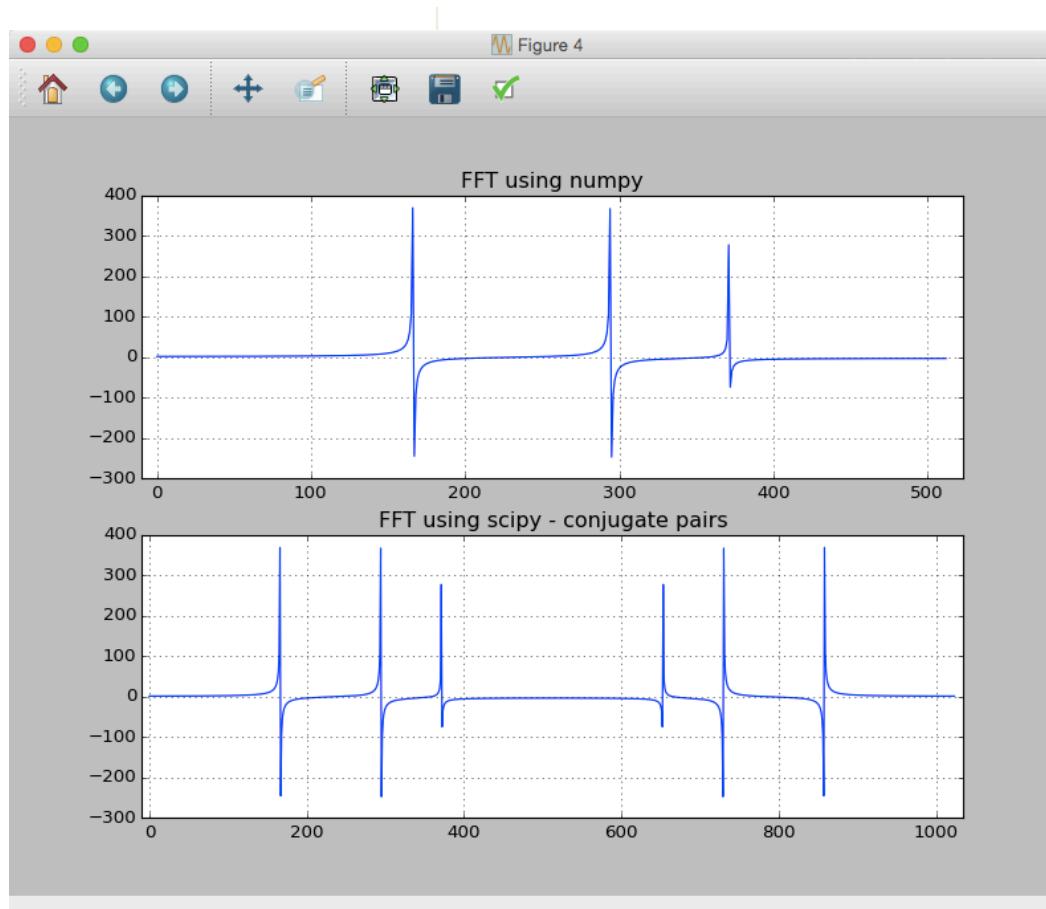
The Fast Fourier Transform

- The Fast Fourier Transform

Example:

```
302 # Plot the Frequency domain signal using numpy:  
303 figure(4)  
304 subplot(2,1,1),  
305 plot(freq_domain_npy),  
306 xlim(-10,len(t)+10),  
307 title('FFT using numpy'),  
308 grid(True)  
309  
310 # Plot the Frequency domain signal using scipy:  
311 subplot(2,1,2),  
312 plot(freq_domain_cpy),  
313 xlim(-10,len(freq_domain_cpy)+10),  
314 title('FFT using scipy - conjugate pairs'),  
315 grid(True)
```

- notice the difference
between:
rfft from NumPy and
fft from Scipy

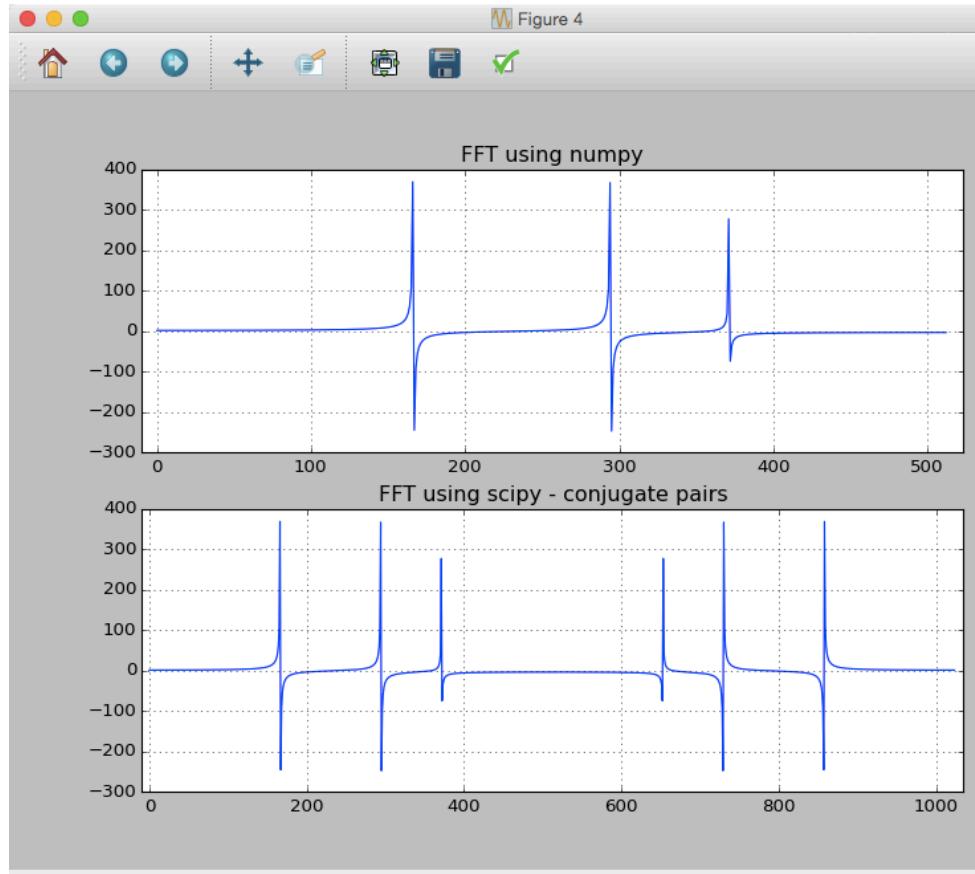


The Fast Fourier Transform

- The Fast Fourier Transform

Example:

- the FFT of a real-valued input signal produces a **conjugate symmetric** result with positive frequencies and their negative counterparts
 - so the negative frequencies are just mirrored duplicates of the positive frequencies
- therefore they **can just be ignored** when analyzing the result
- they are **needed when we want to reconstruct the signal** back to time domain using the inverse FFT the IFFT

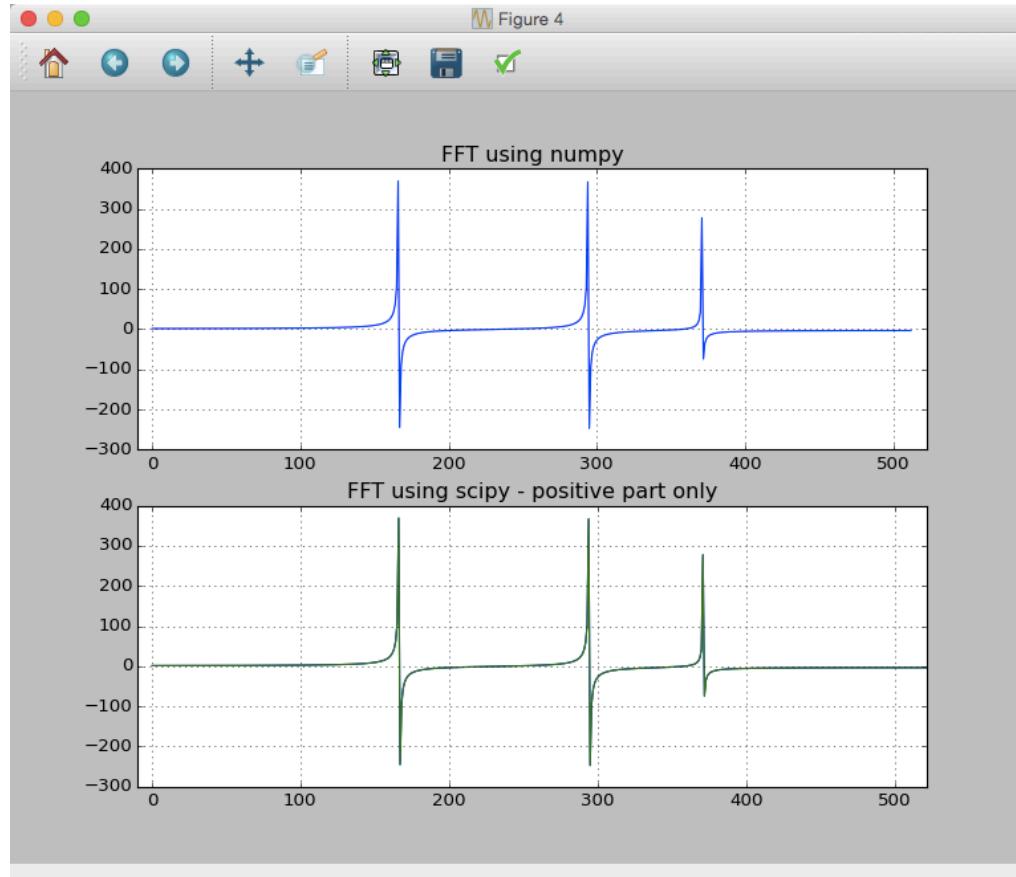


The Fast Fourier Transform

- The Fast Fourier Transform

Example:

- the FFT of a real-valued input signal produces a **conjugate symmetric** result with positive frequencies and their negative counterparts
 - so the negative frequencies are just mirrored duplicates of the positive frequencies
- therefore they **can just be ignored** when analyzing the result
- they are **needed when we want to reconstruct the signal** back to time domain using the inverse FFT the IFFT



Signal Processing

- Signal Processing
 - there are many **Digital Signal Processing** (DSP) techniques we can use and develop using Scipy that can be applied to one-dimensional and multi-dimensional arrays on signals such as **music**, **speech** and **images**
 - many of the signal processing tools are located under the **scipy.signal** package
 - some of the DSP techniques available in Python are:
 - **spectrum analysis** using different windows – Hamming, Hanning, Bratlett, Blackman, Kaiser
 - recreating the time-domain **signal's envelope** – using the Hilbert transform
 - performing **autocorrelation** – correlation of a signal with itself
 - applying **convolution** – convolving two signals (ex: reverb)
 - digital **filter design** – Finite Impulse Response (FIR) filter design
 - using **transforms** – FFT, DCT, DST, Hilbert, etc.

Signal Processing

- Signal Processing

Below are few of the things that we can do:

- **image processing**
 - filtering
 - geometrical transformation of images
 - morphology
 - measurements
- **sound processing (music and speech)**
 - spectrogram
 - filter design
 - time & frequency domain manipulations
- **speech processing**
 - pitch detection
 - pitch change

Signal Processing

- Signal Processing – **image processing:**

- Scipy has several prebuilt functions for image processing
- The module within Scipy that contains these functions is called **ndimage**
- We are not restricted to use image processing tools from one image package only such as **Scipy**, so we can use **misc** image methods (**Scipy**) as well as **Pillow (PIL)** methods

Example:

```
217 ## 3.6 Image processing:  
218 from numpy import flipud # there is also 'fliplr'  
219 from scipy import ndimage, misc  
220 from PIL import Image  
221 from pylab import imshow, figure, subplot, axis, title, pause  
222 import sys  
223  
224 # Load the file:  
225 img = Image.open('files/lecture9/lena.bmp')  
226 img2 = misc.imread('files/lecture9/lena.bmp', flatten=False)  
227 height, width, channels = img2.shape # we save the image dimentions as parameters  
228  
229 # Notice the difference in image object types:  
230 type(img)  
231 type(img2)  
232  
233 # Show either one of the images the same way:  
234 imshow(img2)  
235 pause(1)
```

when “flatten=True” is the same as monochrome image since 3rd channel is gone

In [1]: type(img)
Out[1]: PIL.Image.Image

In [2]: type(img2)
Out[2]: numpy.ndarray

Notice how the two objects are of different types

Signal Processing

- Signal Processing – image processing: image transformation
 - Once loaded the image packages, the user has access to a number of basic image manipulation functionality, such as **image transformation** and **filtering**

Example:

```
233 # Image transformation:  
234 img = img.convert('RGB') # this line will convert our image to RGB  
235 img_bw = img.convert('1') # this line will convert our image to black and white  
236 img_mn = img.convert('L') # this line will convert our image to monochrome  
237 img_rotate = ndimage.rotate(img, 18) # rotate our image at any angle  
238 img_shft_pos = ndimage.shift(img2, (75,125,0)) # print image with [x,y] offset  
239 img_shft_gam = ndimage.shift(img2, (75,125,1.85)) # change the gamma of the image  
240 img_flip = flipud(img2) # this command will flip (not a 180 deg. turn) our image  
241 img_zoom1 = ndimage.zoom(img2, (1.45,1.25,1)) # this command will zoom our image  
242 img_zoom2 = ndimage.zoom(img2, (1.25,1.45,1)) # this command will zoom our image  
243  
244 # Saving the image:  
245 misc.imsave('files/lecture10/lena_mn.bmp', img_mn)
```



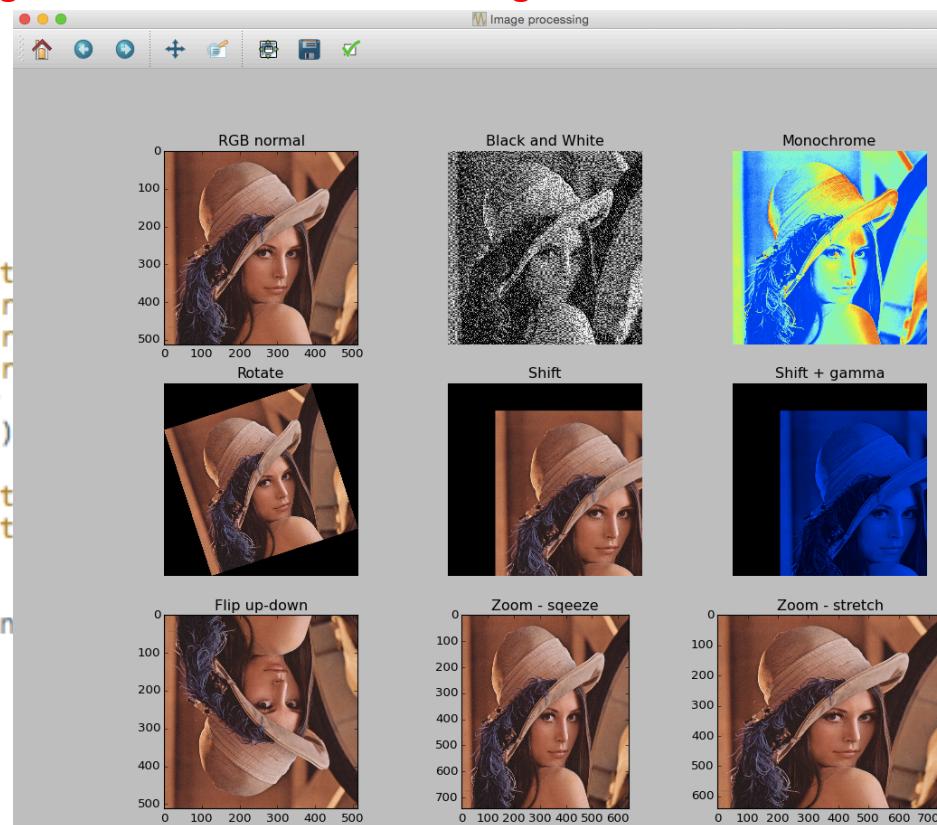
see the saved image

Signal Processing

- Signal Processing – image processing: image transformation
 - Once loaded the image packages, the user has access to a number of basic image manipulation functionality, such as **image transformation** and **filtering**

Example:

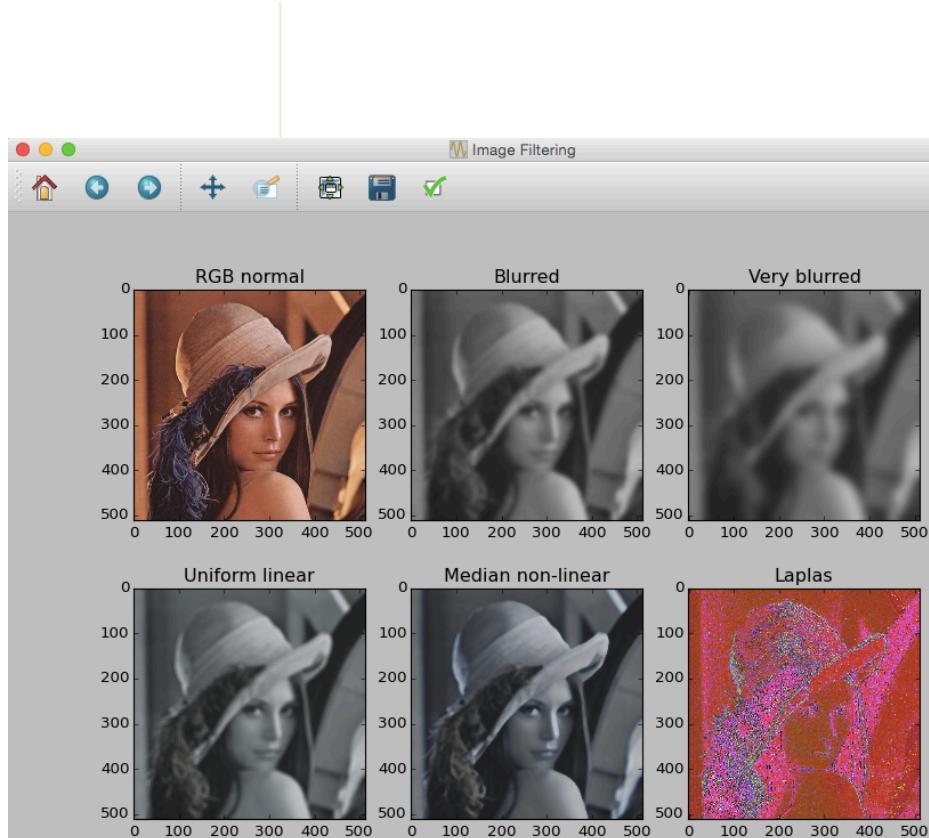
```
233 # Image transformation:  
234 img = img.convert('RGB') # this line will convert  
235 img_bw = img.convert('1') # this line will convert  
236 img_mn = img.convert('L') # this line will convert  
237 img_rotate = ndimage.rotate(img, 18) # rotate our  
238 img_shft_pos = ndimage.shift(img2, (75,125,0)) #  
239 img_shft_gam = ndimage.shift(img2, (75,125,1.85))  
240 img_flip = flipud(img2) # this command will flip  
241 img_zoom1 = ndimage.zoom(img2, (1.45,1.25,1)) # t  
242 img_zoom2 = ndimage.zoom(img2, (1.25,1.45,1)) # t  
243  
244 # Saving the image:  
245 misc.imsave('files/lecture10/lena_mn.bmp', img_mn)
```



Signal Processing

- Signal Processing – image processing: image filtering
 - image

```
259 # Filtering:  
260 from numpy import copy, random, float  
261 from scipy import misc, ndimage, signal  
262 from pylab import imshow, figure, subplot, axis, title  
263  
264 # Load the file:  
265 img = misc.imread('files/lecture10/lena.bmp', flatten=False)  
266  
267 # Gaussian filter:  
268 blurring_1 = ndimage.gaussian_filter(img, sigma=5)  
269 blurring_2 = ndimage.gaussian_filter(img, sigma=10)  
270  
271 # Uniform (linear) filter:  
272 uniform_img = ndimage.uniform_filter(img, size=11)  
273  
274 # Median (non-linear) filter:  
275 median_img = ndimage.median_filter(img, size=9)  
276  
277 # Laplas filter:  
278 laplas_img = ndimage.laplace(img)
```



Signal Processing

- Signal Processing – sound processing: spectrogram
 - spectrogram is a 3-D way of visualizing the frequency domain of any given signal
 - spectrograms are 3-D because they represent frequencies and their magnitudes over time in a given signal
 - signals are usually: sounds, music and speech, but can also be image signals
 - sometimes spectrograms are referred to as waterfalls, voiceprints, or voicegrams
 - they are the perfect tool for phonetic analysis in visualizing spoken words
 - spectrogram visualizing functionality can be uploaded in one of two ways:
 - In [1]: `from pylab import specgram` ... or
 - In [1]: `from matplotlib.pyplot import specgram`
 - sometimes the spectrogram format varies and the vertical and horizontal axes can be switched
 - spectrograms are usually generated in two ways, by using:
 - FFT calculated from a given time signal
 - Filterbanks resulting from a sequence of bandpass filters

Signal Processing

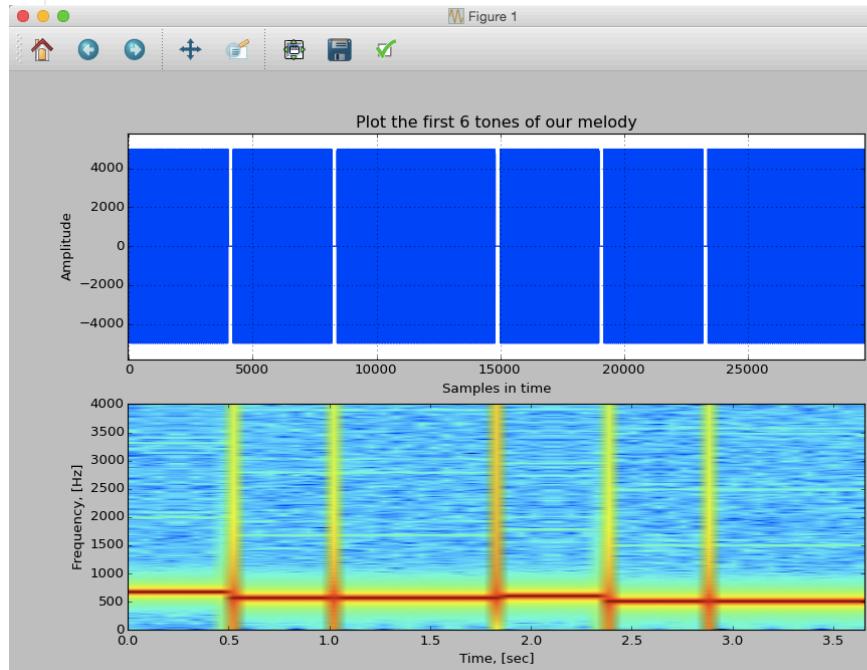
- Signal Processing – sound processing: spectrogram
 - here are some of the parameters users have control over:
 - **Fs** – the **sampling frequency** calculating the Fourier frequencies, in cycles per time
 - **NFFT** – the amount of **frequency bins** represented in each FFT window
 - **window** type used – **Hamming**, **Hanning**, **Bratlett**, **Blackman**, **Kaiser**
 - **noverlap** – amount of **overlap between window blocks**. the default overlap is 128 samples
 - **mode** – type of spectrogram to visualize: { 'psd' | 'magnitude' | 'angle' | 'phase' }, where:
 - **psd** – is the power spectral density
 - **magnitude** – is the magnitude spectrum
 - **angle** – represents the phase spectrum without unwrapping
 - **phase** - is the phase spectrum with unwrapping
 - **scale** – how the data should be displayed { 'default' | 'linear' | 'dB' }, where:
 - **default** – linear
 - **linear** – means no scaling will be used
 - **dB** – when mode='psd' the dB scale is $(10 * \log_{10})$, otherwise it is $(20 * \log_{10})$
 - **Fc** – the center frequency can be controlled
 - **cmap** – is the colormap chosen

Signal Processing

- Signal Processing – sound processing: spectrogram

Example:

```
1 ## Scipy 3/3
2
3 # Signal Processing:
4 # 1. Spectrogram:
5
6 from pylab import specgram, plot, subplot, title, xlabel, ylabel, grid, axis, xlim, ylim
7 from scipy.io.wavfile import read
8
9 (Fs, x) = read('files/lecture10/melody.wav') # Fs - sampling frequency, x - signal
10
11 # Lets plot signal 'x':
12 subplot(2,1,1)
13 plot(x)
14 xlim([-50,29750]); ylim([-5800,5800]) # limit it to the first 6 tones only
15 title('Plot the first 6 tones of our melody')
16 xlabel('Samples in time'); ylabel('Amplitude')
17 grid(True)
18
19 # Now we plot the spectrogram of the first 6 tones:
20 subplot(2,1,2)
21 xlabel('Time, [sec]'); ylabel('Frequency, [Hz]')
22 specgram(x[0:29750], NFFT=512, Fs=Fs, Fc=0, noverlap=12, mode='magnitude')
23 axis('tight')
```



Signal Processing

- Signal Processing – sound processing: spectrogram

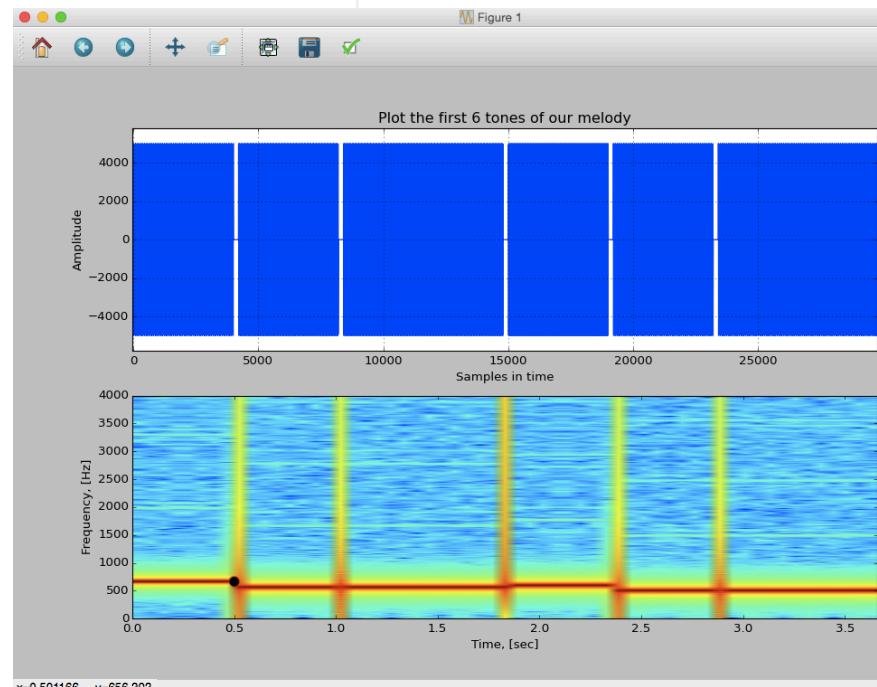
Example:

lets recall the code snipped from the melody:

```
97 # Creating each tone with Fs samples per second and length 0.5 or 0.8 seconds:  
98 Es = note(659,Fs,0.5,amplitude=5000)  
99 El = note(659,Fs,0.8,amplitude=5000)  
100 Css = note(554,Fs,0.5,amplitude=5000)  
101 Csl = note(554,Fs,0.8,amplitude=5000)  
102 D = note(587,Fs,0.5,amplitude=5000)  
103 Bs = note(494,Fs,0.5,amplitude=5000)  
104 Bl = note(494,Fs,0.8,amplitude=5000)  
105 As = note(440,Fs,0.5,amplitude=5000)  
106 Al = note(440,Fs,0.8,amplitude=5000)
```

and the melody: E C# C# D B B A B C# D E E E
E C# C# D B B A C# E E A A A

- notice how the frequency tones are represented by the ‘hottest’ red line in the spectrum
- we also notice how the timing of each tone corresponds to 0.5 and 0.8 sec
- for E: x=0.501166 [sec], y=656.303 [Hz]



Signal Processing

- Signal Processing – sound processing: envelope
 - the envelope of a time-domain signal is a smooth curve around the extreme parts of that signal so it connects all the peaks in a signal and is the **estimate of the total signal level** of the audio input
 - the envelope detection methodology of **squaring and lowpass filtering** a signal works like this:
 - it **squares the input signal**, due to which half of its energy is shifted up to higher frequencies and half is going down toward DC
 - then sending it through a minimum-phase lowpass filter (**LPF**), to eliminate any high frequency energy thus **avoid aliasing** (FIR decimation)
 - the signal is **downsampled**
 - simpler way of calculating the envelope of a signal is to just use the **Hilbert transform**
 - the **Hilbert transform** is used to calculate the **analytic signal** – one **without negative components**. They can be discarded at no loss
 - the analytic signal is a complex signal, with real part represented by the original signal and imaginary part by the Hilbert transform of that same signal
 - the **envelope signal can be used to observe how the signal changed** after some frequency components were suppressed or reinforced by filtering of the original signal

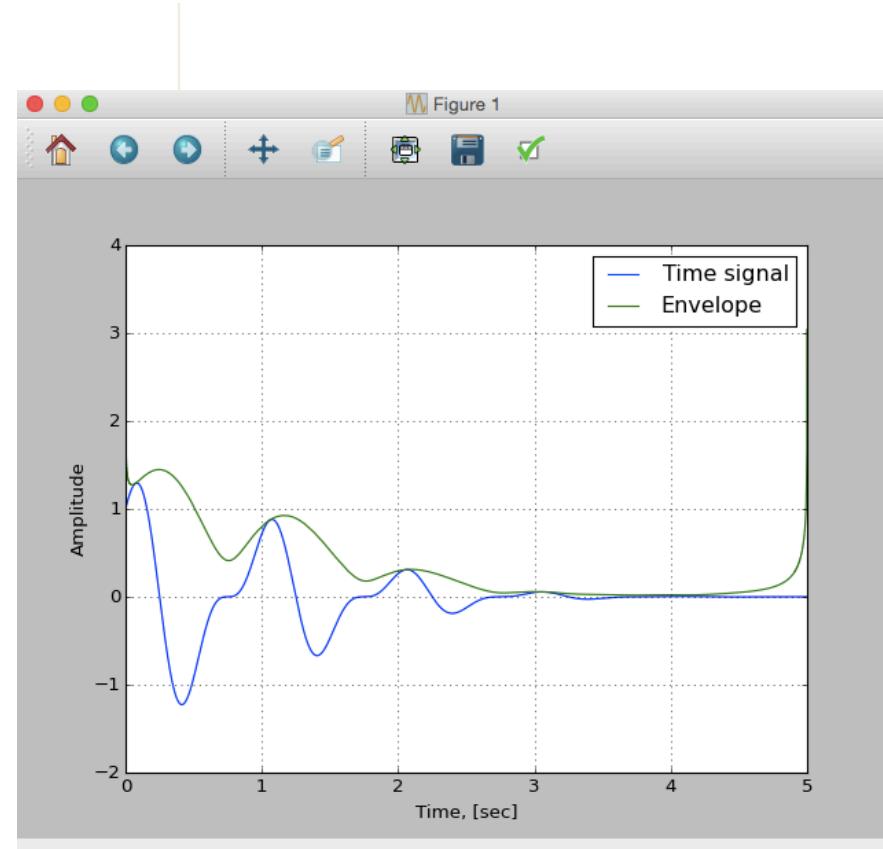
Signal Processing

- Signal Processing – sound processing: envelope

Example:

```
25 # 2. Envelope:  
26 import scipy.signal.signaltools as sigtool  
27 from numpy import linspace, sin, cos, pi, exp, abs  
28 from pylab import plot, grid, legend, show, xlabel, ylabel  
29  
30 # Lets create a simple signal using the 'sin' and 'cos' functions:  
31 a = linspace(0,5,1e5)  
32 b = (1+sin(2*pi*a))*exp(-a**2/3)*(cos(2*pi*a))  
33  
34 # Now we take the envelope using 'Hilbert's' transform:  
35 envelope = abs(sigtool.hilbert(b))  
36  
37 # Plotting:  
38 plot(a,b,label='Time signal')  
39 plot(a,envelope,label='Envelope')  
40 xlabel('Time, [sec]'), ylabel('Amplitude'), legend(), grid(True)  
41 show()
```

- the shape of the envelope matches the actual signal
- the amplitude of the envelope does not match the original signal, because the Hilbert transform was implemented using a non-ideal FIR filter, so the magnitude response is not one for all frequencies



Signal Processing

- Signal Processing – sound processing: speech

Vowels and phonemes
in American – English

We use five letters to
represent the vowel sounds:
a, e, i, o, u

IPA fonts in Ladefoged (2006) and Roach (2009)

Words ^②	Ladefoged (2006) ^③	Roach (2009) ^④	Words ^②	Ladefoged (2006) ^③	Roach (2009) ^④
feet ★ ^⑤	/i/ ^⑥	/ɪ/ ^⑥	bird★ ^⑤	/ɜ:/, /ə:/ ^⑥	/ə/ ^⑥
hard★ ^⑤	/ɑ/ ^⑥	/a:/ ^⑥	bed★ ^⑤	/ɛ/ ^⑥	/e/ ^⑥
food★ ^⑤	/u/ ^⑥	/ʊ/ ^⑥	attend ^⑦	/ə/ ^⑥	/ə/ ^⑥
ord★ ^⑤	/ɔ/ ^⑥	/ɔ:/ ^⑥	book ^⑦	/ʊ/ ^⑥	/ʊ/ ^⑥
hot★ ^⑤	/a/(GA), ^⑧ /ɒ/ (RP) ^⑨	/ɒ/ ^⑩	go★ ^⑤	/oʊ/(GA), ^⑧ /əʊ/(RP), ^⑨	/əʊ/ ^⑩
bus, ^⑩	/ʌ/ ^⑪	/ʌ/ ^⑪	boy ^⑩	/ɔɪ/ ^⑫	/ɔɪ/ ^⑪
book ^⑩	/ʊ/ ^⑪	/ʊ/ ^⑪	big ^⑩	/ɪ/ ^⑫	/ɪ/ ^⑪
dear, ^⑩	/ɪə/ ^⑪	/ɪə/ ^⑪	care★ ^⑤	/ɛə/ ^⑫	/eə/ ^⑪
bike ^⑩	/aɪ/ ^⑪	/aɪ/ ^⑪	how ^⑩	/au/ ^⑫	/au/ ^⑪
cake ^⑩	/eɪ/ ^⑪	/eɪ/ ^⑪	tour ^⑩	/uə/ ^⑫	/uə/ ^⑪

★ =There is a difference between two systems^⑪

Signal Processing

- Signal Processing – sound processing: speech

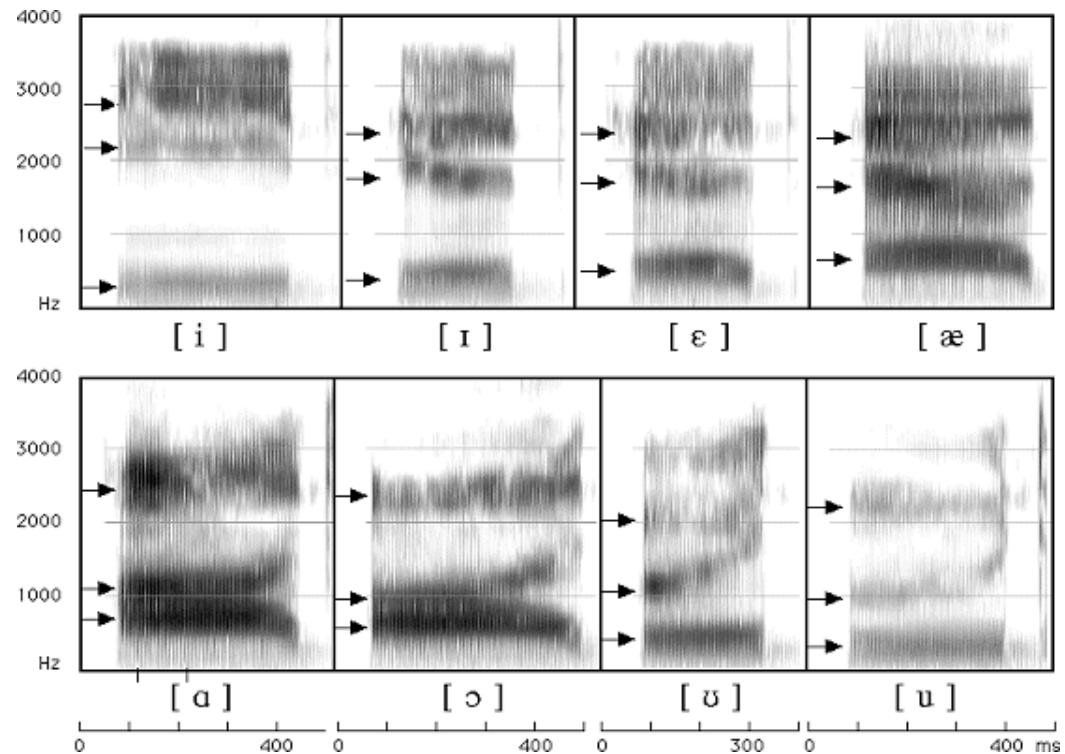
Formants are resonances in the vocal tract

Each vowel is formed by formants: a concentration of acoustic energy around a specific frequency

Each formant has a different center frequency with higher amplitude

Formants for different genders or kids vary

Spectrograms of the American English Vowels



(Ladeforged 2006:185-187)

Signal Processing

- Signal Processing – sound processing: speech

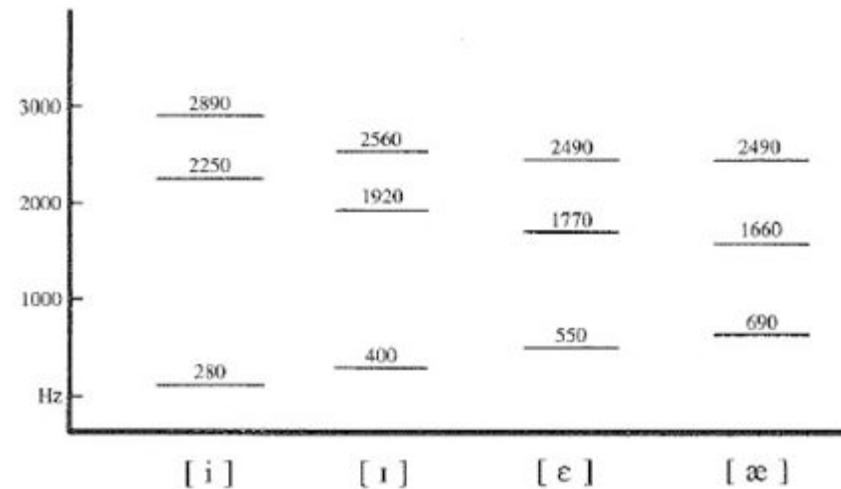
Formants are resonances
in the vocal tract

Each vowel is formed by
formants: a concentration
of acoustic energy around
a specific frequency

Each formant has a
different center frequency
with higher amplitude

Formants for different
genders or kids vary

Spectrograms of the American English Vowels



(Ladefoged & Johnson, 2011:193)

Signal Processing

- Signal Processing – sound processing: speech

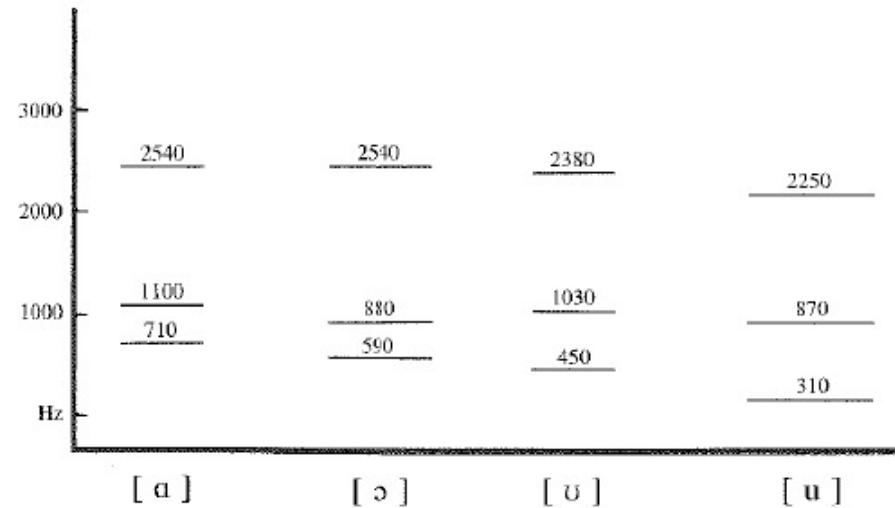
Formants are resonances
in the vocal tract

Each vowel is formed by
formants: a concentration
of acoustic energy around
a specific frequency

Each formant has a
different center frequency
with higher amplitude

Formants for different
genders or kids vary

Spectrograms of the American English Vowels



(Ladefoged & Johnson, 2011:193)

Signal Processing

- Signal Processing – sound processing: speech

We use five letters to represent the vowel sounds:
a, e, i, o, u

use LPC for formant estimation from [Audiolazy](#):

[pip install audiolazy](#)

Formant frequencies for each vowel

Vowel	F1(Hz)	F2(Hz)	F3(Hz)
i:	280	2620	3380
I	360	2220	2960
e	600	2060	2840
æ	800	1760	2500
ʌ	760	1320	2500
a:	740	1180	2640
ɒ	560	920	2560
ɔ:	480	760	2620
ʊ	380	940	2300
u:	320	920	2200
ɜ:	560	1480	2520

Adult male formant frequencies in Hertz collected by J.C. Wells around 1960.
Note how F1 and F2 vary more than F3.

Signal Processing

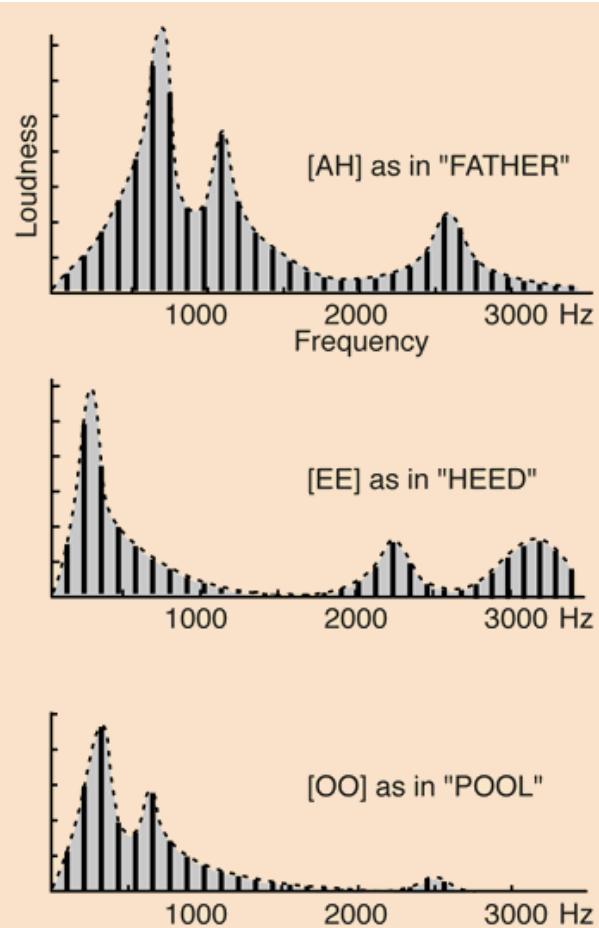
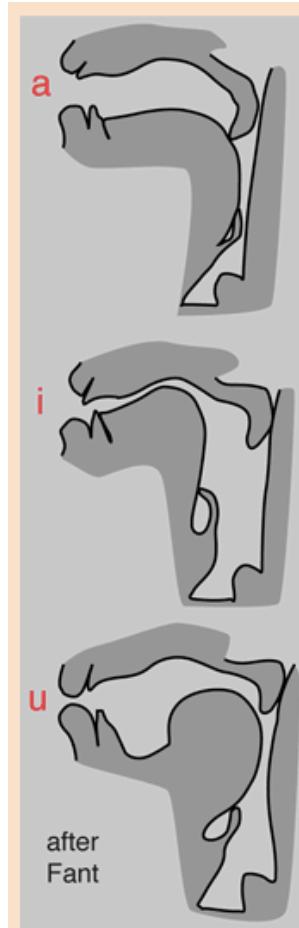
- Signal Processing – sound processing: speech

See levels for each formant

Each vowel is formed by formants

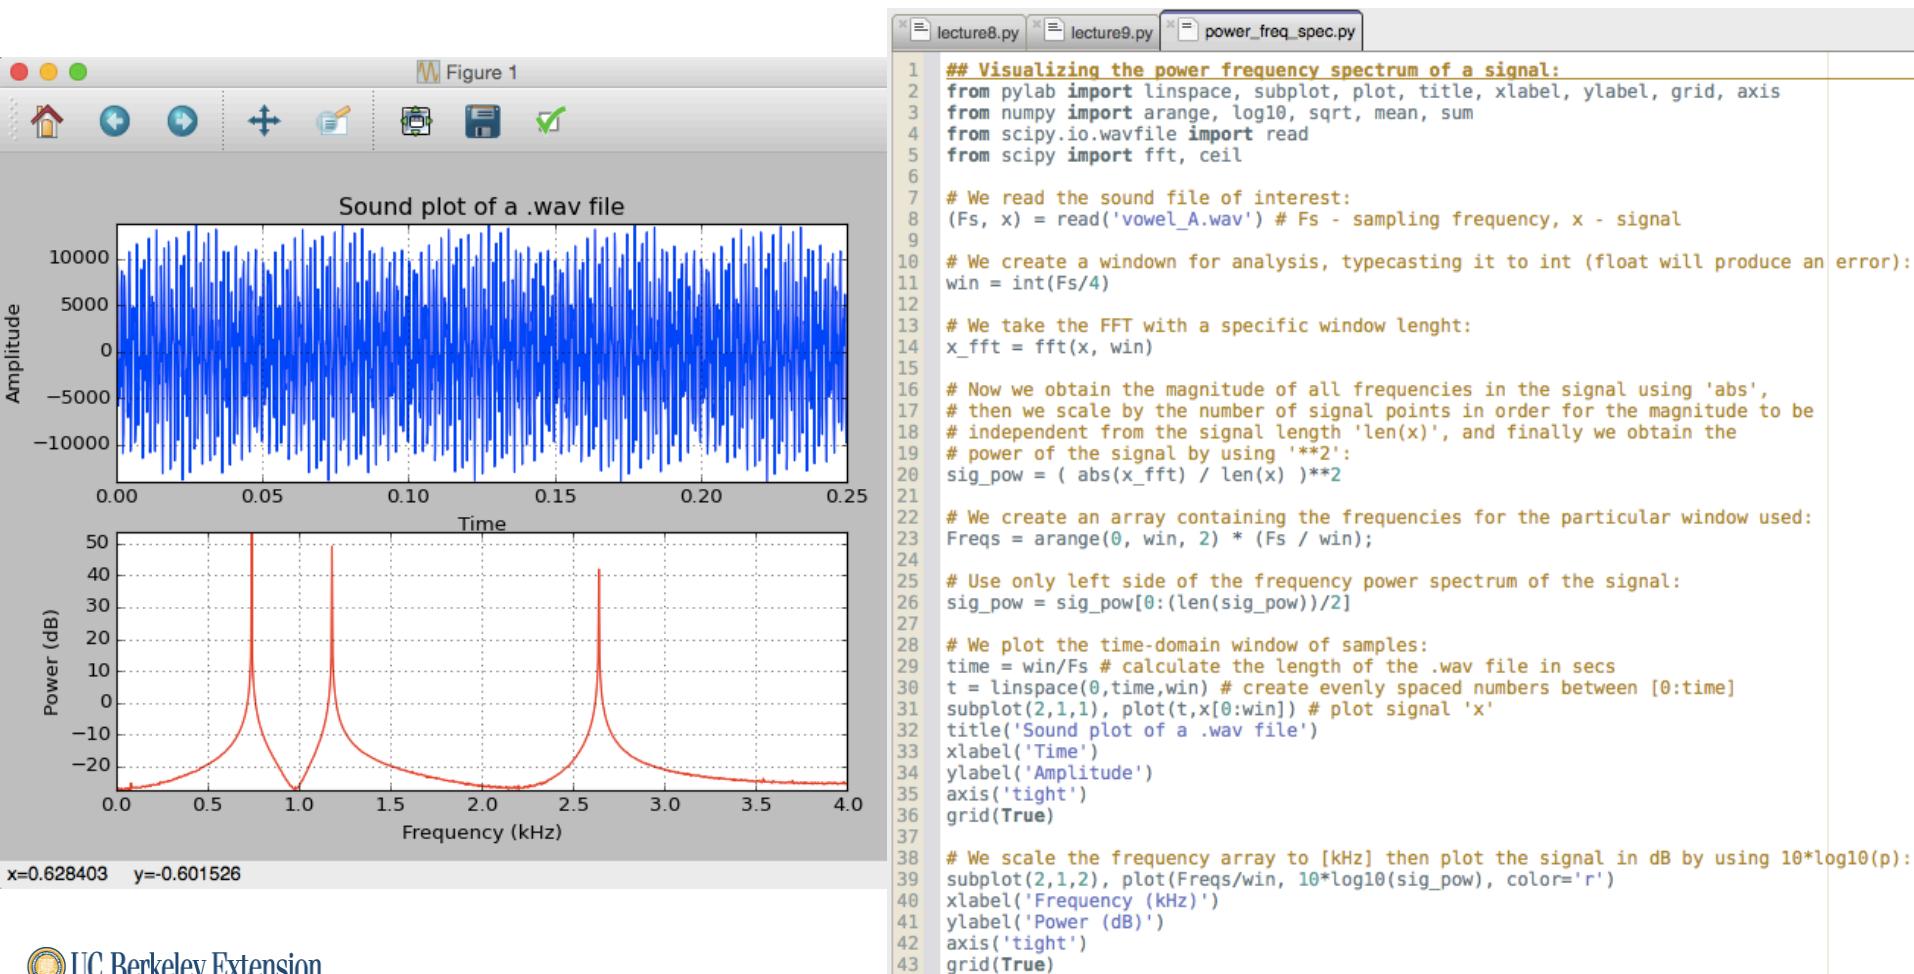
Each formant has a different frequency and amplitude

Formants for different genders or kids vary



Signal Processing

- Signal Processing – sound processing: power frequency spectrum



Homework #4

Generate the English phonemes corresponding to the vowels: a, e, i, o, u

1. Synthesize each formant tone in preparation for creating each vowel

We will simplify the synthesis of each vowel: a, e, i, o, u accepting the following constraints:

1. you will use pure tones only (serving as center frequencies) using the sin function
2. you will manually adjust the amplitude of each formant frequency you generate
3. you will use the first three formants for the average male speaker

Note: in the exercise we will not use: glottal pulse train simulation, vocal tract estimation, female/child speaker formants, formant bandwidth reproduction, so it will be a simplified model.
Only addition of simple tones will be used to represent each vowel.

2. Create a function that generates each simple formant tone with sampling frequency $F_s=8\text{kHz}$, duration 2 seconds, and different amplitudes based on the spectrograms of the American English Vowels provided by Ladeforged 2006:185-187. Refer to the spectrogram provided in the lectures and use your judgment. Obtain best results by testing and repetition!
3. Use the formant frequencies given by J.C. Wells - refer to the table provided in the lectures
4. Simply add all three formants for each vowel and save the files in .wav format
5. **Extra point (optional):** you can incorporate the code I provided in the previous slide as a module.py to plot each vowel sound. In this case send me both files.
6. Organize your code: use each line from this HW as a comment line before coding each step
7. Save these steps in a .py file and email it to me before next class. I will run it!